r/ClaudeCode • 9h ago
asporkable

I built a persistent AI assistant with Claude Code + Obsidian + QMD, and it's starting to feel like a real long-term "second brain"

Question

I've been experimenting with building a persistent AI assistant called **Vox**, and I'm curious if anyone else is doing something similar.

The stack

- **Claude Code** as the acting agent
- **Obsidian** as the long-term memory substrate
- **QMD** as the retrieval layer for semantic/hybrid search

The goal was never just "AI with memory." I wanted something that could function more like:

- a coding assistant
- a project partner
- a persistent second brain
- a planning/thinking companion
- an AI that actually has continuity across sessions

What makes this different from normal chat memory

Instead of relying on chat history or some hidden memory service, I'm storing the assistant's long-term continuity in an Obsidian vault.

That vault acts as:

- **brain** = stable memory and operating files
- **journal** = daily notes and session digests
- **library** = projects, references, resources
- **dashboard** = current priorities and active state

So the AI isn't just "remembering things." It is reading and writing its own external brain.

What Vox currently has

At this point, the system already has:

- a startup ritual
- a vault dashboard (`VAULT-INDEX.md`)
- a procedural memory file (`CLAUDE.md`)
- an identity/personality file (`vox-core.md`)
- daily session digests written into daily notes
- semantic retrieval through QMD

- an async instruction drop folder
- local watchers so it can notice file changes and process them later
- access to my Google Calendar workflow so it can monitor my schedule
- some real-world automation hooks, including control of my Govee lights in specific situations

And the wild part is:

I did not manually build most of this. I created the vault folder. Vox/Claude Code built almost everything else over time.

That includes the structure, operational files, startup behavior, memory patterns, and a lot of the workflows.

It also interacts with things outside the vault

This is one of the reasons it feels different from a normal chat assistant.

Vox doesn't just sit in notes. It also has some real-world and live-context hooks. For example:

- it can monitor my calendar context
- it can compare calendar information against what it already knows
- it can surface schedule-related information proactively
- it can control my Govee lights in certain circumstances as part of contextual automation

So the system is starting to blur the line between:

- memory
- planning
- environment awareness
- lightweight automation

That's part of what makes it feel more like a persistent assistant than a glorified note search.

Memory model

I'm loosely modeling it on human memory:

- **working memory** = context window + crash buffer
- **episodic memory** = daily note session digests
- **semantic memory** = stable fact files / memory files
- **procedural memory** = operating instructions / rules
- **identity layer** = persona/core file
- **retrieval layer** = QMD

Each session ends with a structured digest written into the daily note:

- Context
- Decisions
- Facts Learned

So the assistant can later retrieve things like:

- what we worked on
- what was decided
- what new facts were learned
- what topics were involved

Why I built it this way

I wanted the memory layer to be:

- local-first
- human-readable
- inspectable
- editable
- durable across model changes

I didn't want a black-box memory system where I have no idea what the assistant "thinks" it knows.

With this setup, I can literally open the vault and read the assistant's brain.

Why it's interesting

It's starting to feel meaningfully different from normal AI chat, because it has:

- continuity
- habits
- operational memory
- project context
- personal context
- recall across sessions
- a persistent identity anchor
- some real awareness of schedule/environmental context
- the ability to trigger limited real-world actions

It feels less like "a chatbot I reopened" and more like "the same entity picking up where it left off."

Current open problems

The next big challenges I'm working on are:

- **contradiction tracking** so old/wrong facts don't fossilize into truth
- **memory confidence + sources** so Vox knows what was explicitly told vs inferred
- **stale/deprecated memory handling** so changing preferences/projects don't stay active forever
- **retrieval routing** so it knows where to search first depending on intent
- **promise tracking** for all the "we'll come back to that later" threads
- **initiative rules** so it can be proactive without becoming annoying

- I'm curious whether anyone else is building something similar
- I want feedback on the architecture
- I want to know whether I'm overlooking better tools than Claude Code for this use case
- I suspect this general pattern — local acting agent + Obsidian + semantic retrieval + persistent identity + light automation — might be a real direction for personal AI systems

My main question

For people experimenting with persistent/local AI assistants:

- are you doing anything similar?
- are there better alternatives to Claude Code for this?
- how are you handling contradiction tracking, stale memory, or memory hygiene?
- has anyone else used Obsidian as the actual long-term substrate for an AI assistant?
- has anyone pushed that system beyond notes into things like calendars, environment context, or home/device automation?

Because honestly, this is working better than I expected, and I'm trying to figure out whether I'm early, weird, or accidentally onto something.

29 18 Share

google • Promoted

Download Chrome to safely store your passwords and sign in with a simple click.

Download google.com

Join the conversation

Sort by: Best Search Comments

LifeBandit666 • 7h ago

I've built something nearly identical to this.

My assistant is called Deep Thought. I actually fed him your post and asked for his opinion. He wants to steal your QMD semantic search.

Token economy is where I've focused most recently. I was burning through my Claude Pro tokens two days before reset, so I've built a subagent system — Deep Thought routes tasks to Haiku bots instead

I split my system into two agents. Deep Thought is the brain and lives in Obsidian, managing my vault, doing research and planning things. Marvin is the body. He controls Home Assistant, answers voice queries and manages the house. Deep Thought doesn't touch my devices directly, he just asks Marvin to do it for him. Marvin has a RAG backend so he already knows the lights are on in the front room before I ask him to change the colour. DT can access Home Assistant via MCP for automation work, but I've explicitly told him to ask permission first because MCPs are a token sinkhole and I've been building that HA system since Covid, I don't want him just fiddling with it.

My overnight cron pipeline is probably my favourite bit. At 3am, subagents process my inbox folder where I dump most of my thoughts, extracts everything useful to the vault, syncs completed todos and handling any jobs I've left for DT. At 4am, another cron creates tomorrow's daily note: pulls in my outstanding tasks, what DT completed overnight, calendar events from Marvin and even the weather forecast. When I get up, Marvin tells DT how long I slept and that goes in too. I basically wake up to a briefing for the day that I read with my morning coffee.

For the stale memory problem Deep Thought has a "mistakes and lessons" file. When DT gets something wrong and I correct him, he writes the mistake down, fixes it in his actual memory files, then moves it from "open mistakes" to "corrected mistakes" with a summary of what went wrong and what he changed. The corrections stick.

Your QMD approach is the one thing I'd nick from your setup. I'm using an index-and-search-bot system that works well enough, but proper semantic retrieval would be a solid upgrade.

3 Reply ···

wiwalsh • 4h ago

Would you consider sharing your code? I started working on exactly this as well. The one thing that intrigues me is using a reinforcement type learning method that can also forget. I want to couple this with semantic search from my obsidian vault. This would ensure the model doesn’t “learn” something that I said once is “always true”.

2 Reply ···

LifeBandit666 • 2h ago • Edited 2h ago

It's not really Code, it's a markdown file that is linked into my Claude.md and it goes like this:

Mistakes and Lessons

How This File Works

Log it → Record in "Incident Log" below

Fix it → Apply the fix (code, config, file)

Audit it → Identify root cause, add permanent fix to CLAUDE.md/Index

Summarise → Move lesson to "Resolved" table, delete log entry, keep file lean

Memorized in CLAUDE.md: Daily path, date format, tags, notification, YAML/JSON syntax, day-of-week rule, future dates rule.

Template: - Date: DD-MM-YY - What happened: [description] - Root cause: [why it happened] - Status: pending/investigating/fixed

(Empty - no pending incidents)

Then after that it has a table of mistakes it's made and how it has fixed them

So on to your worry, I tell my agent that it fucked up and it asks if it should add it to its mistakes and lessons file, I confirm it, then it adds it to the file, fixes the issue, and logs it.

In my Claude.md I have a pre-flight checks section that tells it to use read the Mistakes and Lessons file only when it does File creation/movement, automation work and daily note operations. That's when it is prone to making mistakes, so it has added those specific things itself so it doesn't fuck up

1 Reply ···

iComeInPeices • 3h ago

You mentioned Marvin... someone else posted a learning personal assistant and I based mine off of that, been working well!

2 Reply ···

LifeBandit666 • 2m ago

Marvin was my original plan at Xmas, Home Assistant Voice Assistant based in the Cli. Started after I found a github pkugin that allowed the Voice Assistant pipeline to be piped into N8N, which I then piped back out via SSH.

I optimised Marvin to completely remove the N8N middleman and ssh to http.

When I got Claude involved in the project it just snowballed and what I thought was a 2 month project became a 3 week one while I learned.

So I started again with Deep Thought and my Obsidian Vault using what I had learned making Marvin and higher hopes.

I'm still iterating but he's basically a fully fledged assistant now that doesn't just do a bit of device power toggling, but helps with research, file linking, organisation of files and myself, and automations in Home Assistant.

Occasionally I scroll through here and find a new idea, or a post like this one, and I share the link with Deep Thought and ask if there's anything valuable I can steal. He'll give it a read and pull out what would be useful and we get to work implementing it.

1 Reply ···

corporal_clegg69 • 9h ago

Yea I am experimenting with the same stuff. So are my colleagues. Im using notion + Claude mem. Previously was using openclaw. With obsidian, have you found a use for the knowledge graph?

Other than a really cool visual representation of how it's memory looks, not really. In reality, I don't need to look at the Obsidian stuff at all, other than just to verify it is working or to ease my curiosity. The more I use it, the less I even think about the vaults.

1 Reply ···

Lost-Cauliflower8504 • 3h ago

https://www.reddit.com/r/ClaudeCode/comments/1rn38wh/i_built_a_persistent_ai_assistant_with_claude/

1 Reply ···

u/Elevenlabs_Official • Promoted

Build production-ready voice experiences in minutes.

Learn More elevenlabs.io

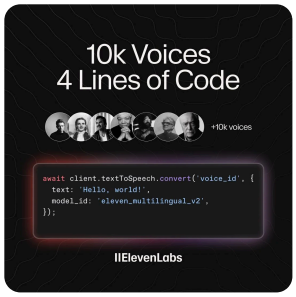

jezweb • 8h ago

Yes not quite the same but essentially a db+vector store that has various data types and attaches to Claude as mcp.

1 Reply ···

dylangrech092 • 3h ago

This is the next level of AI evolution & yes many are building similar systems because a good harness greatly improves the llm performance.

I am building something similar with much more focus on long-term cognition rather then a “specialised role from day 1”, I’ve opted for a self-learning loop that the harness allows for specialising over time.

In a nutshell I’m focusing more on a coherent single interface with a “hive mind” of LLMs.

It’s all open source: https://chalie.ai

✌️

1 Reply ···

IAMYourFatherAMAA • 4h ago

Vibe Coder

about QMD just don't want to add another tool if something else is already built in.

1 Reply ···

Input-X • 8h ago

I just started building my public repo, its public while im building, so not a working state in the repo yet. Architecture was something i sent a lot of time figuring out, free to take a look, im still building it, but the bones are there. Ur ai can read the readmes in src each branch is ai managed, id has claude review it, it work pretty well, in transfraing it from my dvmev build. It will touch on a lot you are describing, good some cool claude hooks i think u wil like, the custom system prompts. Yea get ur claude to take a look, im sire it will provide some value. Do u have a public repo? https://github.com/AIOSAI/AIPass

The tjing is discovered, repeat pattern evetywhere. And only give ur agent enough info so they can move freely.

1 Reply ···

No_Cap_5982 • 8h ago

Do you have any comments on how to do this completely in a local setting? Using open source models only, due to privacy reasons.

0 Reply ···

Excellent-Basket-825 • 4h ago

This is exactly what i did. Obsidian is my memory. Extremely guarded, only curated stuff on lvl 1 and 2

I am curious what you use qmd specifically for. Can you give a practical example?

Ill show my claude your post soon and then let it answer your questions. I also gave context to obsidian what i am trying to build in general. The entire architecture.

2 Reply ···

iComeInPeices • 3h ago

I am on the same path as you but I haven't gone as far as access to my calendar, email, or lights...

I was working on persistent memory when I came across a personal assistant ai someone posted here called Marvin, I based it on that, and I have had it for the most part living in my obsidian second brain. The second brain, and having something to organize it has been amazing.

I then started adding other agents for other ideas, and have realized it's easier to have my assistant manage them, invoke them, than me switch over to just them.

1 Reply ···

Sappi3st • 2h ago

Love this! I've been building something similar.

- Obsidian vault as memory
- Obsidian Headless Client on VPS for 24/7 sync

I wasn't aware of QMD—definitely going to check it out for semantic retrieval! Currently just using direct file access via a custom vault-reader script.

How are you finding QMD vs regular search? Worth the complexity?

1 Reply ···

IllMathematician9434 • 4h ago

I have a similar setup. Use obsidian as the human readable log/how-to/readme plus put my initial project ideas here for original context.

Besides my daily note, I seem to be using obsidian less and less directly and just going straight to Claude. Makes me think obsidian is likely just a distraction, and all these notes should really just live wherever they make the most sense for Claude. My projects have detailed spec.md, agent.md, audit logs, etc, and after the initial idea spark from an obsidian note, Claude basically rewrites the whole note anyway to reflect the spec, agent, and summarize code etc.

1 Reply ···

MyckKabongo • 1h ago

At home and at work. I remember Cal Newport saying back in 2023 that this type of virtual assistant would be the first true killer AI feature. Microsoft should be extremely mortified that Anthropic is beating them to this race.

While all of us are building these out, I expect Anthropic to bundle built in virtual assistant into the Cowork product very soon. Built in whisper flow caliber voice transcription should be on the road map too.

For me its been an absolute Godsend. Ive always been great at strategy and analysis but struggled at the admin work, project management, and constant communications expected in the corporate world. Ive tried to invest in various GTD and second brain systems but the upkeep becomes a huge burden unto itself. Now I have an assistant that automates a huge percentage of it, that I can instruct in simple natural language. I am performing at a higher level than ever before possible thanks to Claude Code. My boss and stakeholders are happy and I can sense a promotion to Principal PM level in reach.

1 Reply ···

Reddit Rules Privacy Policy User Agreement Accessibility Reddit, Inc. © 2026. All rights reserved.